

# CSCI 3485 Lab 4

Rafael Almeida

October 22, 2024

## 1 Introduction & Methodology

This lab reports on the training time and accuracy of four pre-trained models on the full CIFAR10 dataset. The pre-trained models chosen were ResNet18, ResNet34, VGG11, and VGG13.

For transfer learning, only the models' final linear layer was updated to have its *out\_features* equal to 10 (the number of classes in CIFAR10) instead of 1000 (the original number of classes). Nothing else was changed in any of the models. Moreover, the following are the parameters for training all the models:

- Epochs: 5
- Learning Rate: 0.001
- Optimizer: Adam
- Loss Function: Cross Entropy Loss
- Data Transformations: only the transformations provided by each model's definition in PyTorch were used
- Batch Size: 32
- Refer to the source code below for any more information

The results include the total number of parameters and the number of trainable parameters, the accuracy of the models ( $\frac{\text{correct\_classifications}}{\text{total\_data\_points}}$ ), the training time (the execution time of the *Model.train* method), and the predicted classes of two airplane, frog, and horse images arbitrarily chosen from Google Images.

## 2 Results

### 2.1 Main Statistics

Table 1: ResNet18

Total Parameters	11,181,642
Trainable Parameters	5,130
Training Time (s)	620.77
Accuracy	0.7724

Table 2: ResNet34

Total Parameters	21,289,802
Trainable Parameters	5,130
Training Time (s)	700.85
Accuracy	0.7881

Table 3: VGG11

Total Parameters	128,807,306
Trainable Parameters	40,970
Training Time (s)	819.61
Accuracy	0.8212

Table 4: VGG13

Total Parameters	128,991,818
Trainable Parameters	40,970
Training Time (s)	976.45
Accuracy	0.8031

## 2.2 Google Images

Table 5: Classification Results

Image	ResNet18	ResNet34	VGG11	VGG13
Airplane 1	airplane	airplane	airplane	airplane
Airplane 2	airplane	airplane	airplane	airplane
Frog 1	dog	frog	frog	frog
Frog 2	frog	frog	frog	frog
Horse 1	bird	frog	deer	horse
Horse 2	dog	deer	deer	deer



Figure 1: Airplane 1



Figure 2: Airplane 2



Figure 3: Frog 1



Figure 4: Frog 2



Figure 5: Horse 1



Figure 6: Horse 2

### 3 Discussion & Conclusion

ResNet performed considerably worse than VGG, even though it is usually considered more performant than the VGG architectures because of its increased depth. This depth comes with better capabilities in complex tasks. However, the structure of this experiment may have been particularly disadvantageous to the ResNet models because their final linear layer has only 5,130 parameters, which were the only ones that were trained. Despite being shallower, the VGG models may have performed better because their final linear layer is much more significant with 40,970 trained parameters. This vast difference in the number of parameters trained in this experiment may cause the VGG architectures to perform better.

Interestingly, when comparing the models to each other, from ResNet18 to 34, there was some improvement, while there was a worsening performance from VGG11 to 13. This may have occurred because ResNet34 has nearly double the number of parameters to 18, which gives it an advantage. However, from VGG11 to 13, there is not a massive change in the number of parameters, but rather, there is an increased depth in the convolutional layers. It then seems reasonable to conclude that this increased depth in convolutions only becomes beneficial with slightly more epochs. It could be that the increased complexity of VGG13 could not be captured in the 40,970 parameters trained because more training time is required for all the deep features learned by each convolution to show up at the final layer trained.

The models were good at classifying the airplane and decent at classifying the frogs. Maybe the frog images I chose excited the cuteness neurons associated with dogs! Furthermore, the models were terrible at classifying the horses, most of the time getting it confused with a deer. This is probably because the images are in very natural environments and not in a farm environment enclosed by fences, for example. The other noticeable thing is that horses and deer have very similar structures. Probably more training time would have been required for the trained parameters to learn to differentiate between those nuances.

## 4 Code

GitHub Repository

### 4.1 main.py

```
1  """
2  Defines the experiments to be run.
3  """
4
5  from datetime import datetime
6  from time import time
7
8  from torch import cuda, save
9  from torch.backends import mps
10 from torchsummary import summary
11 from torchvision.models import (
12     ResNet18_Weights,
13     ResNet34_Weights,
14     VGG11_Weights,
15     VGG13_Weights,
16     resnet18,
17     resnet34,
18     vgg11,
19     vgg13,
20 )
21
22 from data import append_to_file, classify_images
23 from model import Model
24
25 if cuda.is_available():
26     device = "cuda"
27 elif mps.is_available():
28     device = "mps"
29 else:
30     device = "cpu"
31
32 print(f"Using device: {device}")
33
34
35 def time_it(f):
36     """
37     Times the execution of a function.
38     """
39
40     def wrapper(*args, **kwargs):
41         start = time()
42         result = f(*args, **kwargs)
43         return (time() - start, result)
```

```

44     return wrapper
45
46
47
48 MODELS = {
49     "resnet18": resnet18,
50     "resnet34": resnet34,
51     "vgg11": vgg11,
52     "vgg13": vgg13,
53 }
54
55 WEIGHTS = {
56     "resnet18": ResNet18_Weights,
57     "resnet34": ResNet34_Weights,
58     "vgg11": VGG11_Weights,
59     "vgg13": VGG13_Weights,
60 }
61
62 for name in MODELS:
63     print("my name", name)
64
65     model = Model(
66         MODELS[name],
67         WEIGHTS[name].IMAGENET1K_V1,
68         classes=10,
69         batch_size=32,
70         data_size=-1,
71     )
72     training_time, info = time_it(model.train(
73         device=device, max_epochs=5, min_delta=0, lr=1e-3
74     ))
75
76     # save(model.model, f"./models/{name}.pth")
77
78     accuracy = model.test(device=device)
79
80     file_name = "results.txt"
81     append_to_file(file_name, datetime.now())
82     append_to_file(file_name, f"{name}")
83     append_to_file(file_name, f"param count:
84         {model.get_parameter_count()}")
85     append_to_file(file_name, f"time: {training_time}")
86     append_to_file(file_name, f"info: {info}")
87     append_to_file(file_name, f"accuracy: {accuracy}")
88
89     classify_images(
90         model.model,
91         WEIGHTS[name].IMAGENET1K_V1.transforms(),
92         device,
93         "./images",

```

93        )

## 4.2 model.py

```
1  """
2  Defines a helper Model class to manage the setup, training, and testing
3  of all models.
4  """
5
6  import torch.nn as nn
7  from torch import no_grad, tensor
8  from torch.optim import Adam
9
10 from data import get_data_loaders
11
12 class Model:
13     def __init__(
14         self,
15         model,
16         weights,
17         classes: int,
18         batch_size: int = 32,
19         data_size: int = -1,
20     ) -> None:
21         self.model = None
22         self.train_loader = None
23         self.test_loader = None
24
25         self.weights = weights
26         self.classes = classes
27         self.batch_size = batch_size
28         self.data_size = data_size
29         self.set_model_and_data(model)
30
31     def set_model_and_data(self, model) -> None:
32         """Sets the last layer of the model and the data with the
33         correction transformation"""
34
35         self.model = model(weights=self.weights)
36
37         # freeze all parameters
38         for param in self.model.parameters():
39             param.requires_grad = False
40
41         # change last layer
42         if hasattr(self.model, "classifier"):
43             in_features = self.model.classifier[-1].in_features
```

```

43         self.model.classifier[-1] = nn.Linear(in_features,
44         self.classes)
45
46         elif hasattr(self.model, "fc"):
47             in_features = self.model.fc.in_features
48             self.model.fc = nn.Linear(in_features, self.classes)
49
50         else:
51             raise ValueError("Model architecture not supported")
52
53         self.train_loader, self.test_loader = get_data_loaders(
54             transforms=[self.weights.transforms()],
55             batch_size=self.batch_size,
56             size=self.data_size,
57         )
58
59     def train(
60         self,
61         device: str,
62         max_epochs: int = None,
63         lr: float = 1e-3,
64         min_delta: float = 1e-2,
65     ) -> dict:
66         """
67         Trains the model following the given parameters.
68         """
69
70         self.model.to(device)
71         self.model.train()
72
73         optimizer = Adam(self.model.parameters(), lr=lr)
74         loss_fn = nn.CrossEntropyLoss()
75
76         # info to return about the training
77         # could include the loss/accuracy per epoch for example
78         epochs = 0
79
80         best_loss = float("inf")
81         while max_epochs is None or epochs < max_epochs:
82             epoch_loss = 0
83             # training on each batch
84             for x, y in self.train_loader:
85                 x, y = x.to(device), y.to(device)
86
87                 # avoids error:
88                 # RuntimeError: element 0 of tensors does not require
89                 grad and does not have a grad_fn
90
91                 optimizer.zero_grad()
92                 outputs = self.model(x)

```



```

91         loss = loss_fn(outputs, y)
92         loss.backward()
93         optimizer.step()
94
95         epoch_loss += loss.item()
96
97         epochs += 1
98         avg_loss = epoch_loss / len(self.train_loader)
99
100        # stop if the loss has not improved enough
101        if abs(best_loss - avg_loss) < min_delta:
102            break
103
104        best_loss = min(best_loss, avg_loss)
105
106        return {"epochs": epochs}
107
108    @no_grad()
109    def test(self, device: str) -> float:
110        """
111        Tests the accuracy of the model.
112        """
113
114        self.model.to(device)
115        self.model.eval()
116        # maintain the calculation in the gpu until returning
117        correct = tensor(0, device=device)
118        total = 0
119
120        for x, y in self.test_loader:
121            x, y = x.to(device), y.to(device)
122            predicted = self.model(x)
123            correct += (predicted.argmax(dim=1) == y).sum()
124            total += y.size(0)
125
126        accuracy = (correct / total).item()
127        return accuracy
128
129    def get_parameter_count(self) -> int:
130        """
131        Custom function that returns the number of total and trainable
132        parameters in a model like torchvision's summary function
133        """
134        total_params = sum(p.numel() for p in self.model.parameters())
135        trainable_params = sum(
136            p.numel() for p in self.model.parameters() if p.requires_grad
137        )
138        return total_params, trainable_params

```

### 4.3 data.py

```
1  """
2  Manages the data for the experiments.
3  """
4
5  from os import listdir
6
7  from torch import Tensor, no_grad
8  from torch.utils.data import DataLoader, Dataset, Subset
9  from torchvision.datasets import CIFAR10
10 from torchvision.io import ImageReadMode, read_image
11 from torchvision.transforms import Compose
12
13
14 class CIFAR10Dataset(Dataset):
15     def __init__(
16         self,
17         root: str,
18         train: bool = True,
19         transform: list = [],
20         size: int = -1,
21     ) -> None:
22         self.dataset = CIFAR10(
23             root=root, train=train, download=True,
24             transform=Compose(transform)
25         )
26
27         if size != -1:
28             indices = range(size)
29             self.dataset = Subset(self.dataset, indices)
30
31     def __getitem__(self, i) -> tuple[Tensor, Tensor]:
32         image, label = self.dataset[i]
33         return image, label
34
35     def __len__(self) -> int:
36         return len(self.dataset)
37
38 def get_data_loaders(
39     root: str = "./data/CIFAR10",
40     transforms: list = [],
41     batch_size: int = 32,
42     size: int = -1,
43 ) -> tuple[DataLoader]:
44     test_dataset = CIFAR10Dataset(
45         root=root, train=False, transform=transforms, size=size
46     )
47     train_dataset = CIFAR10Dataset(
```

```

48     root=root, train=True, transform=transforms, size=size
49 )
50
51 train_loader = DataLoader(
52     train_dataset, batch_size=batch_size, shuffle=True
53 )
54 test_loader = DataLoader(
55     test_dataset, batch_size=batch_size, shuffle=False
56 )
57
58 return train_loader, test_loader
59
60
61 def append_to_file(filename: str, data: any) -> None:
62     with open(filename, "a") as file:
63         file.write(str(data) + "\n")
64
65
66 def classify_images(model, transforms, device, path: str):
67     """
68     Classifies images in a directory using a trained model.
69     """
70     classes = [
71         "airplane",
72         "automobile",
73         "bird",
74         "cat",
75         "deer",
76         "dog",
77         "frog",
78         "horse",
79         "ship",
80         "truck",
81     ]
82
83     model.eval()
84     model.to(device)
85
86     for image_name in.listdir(path):
87         image_path = path + "/" + image_name
88         # Read image in RGB mode
89         image = read_image(image_path, mode=ImageReadMode.RGB)
90         image = transforms(image)
91         image = image.unsqueeze(0)
92         image = image.to(device)
93
94         with no_grad():
95             output = model(image)
96
97         append_to_file(

```

```
98         "results.txt", f"{image_path} {classes[output.argmax()]}"
99     )
```